

SATHWIK H NAIK

| (240) 438-0411 | sathwikhnaik@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#) |

PROFESSIONAL SUMMARY

Analytics engineer with hands-on SQL, dbt, and data-modeling experience at a Fortune 50 healthcare organization. Completed M.S. in Data Science at UMD in 2026, backed by self-directed dbt and Snowflake-based analytics engineering capstones.

EDUCATION

University of Maryland, College Park M.S. in Data Science Coursework Statistical Inference, Deep Learning, Data Modeling, Data Visualization, Experimental Design, Advanced Time Series	2024 – 2026 GPA 3.81/4.0
Reva University, India B.Tech, Computer Science and Engineering	2019 – 2023 GPA 3.55/4.0

WORK EXPERIENCE

Elevance Health Software Engineer	Aug 2023 – Aug 2024
- Documented report logic, metric definitions, and validation rules enabling governed self-service reporting - Designed reusable KPI structures, translating operational requirements into consistent, self-serve cross-team metrics - Developed ETL workflows to source, clean, and validate healthcare data, improving quality for stakeholder reports - Automated recurring data-preparation and reporting workflows, eliminating 20+ hours of manual weekly effort - Built and maintained 10+ Power BI dashboards with DAX measures for multiple internal teams	

PROJECTS

Healthcare Insurance Claims Analytics & Fraud Signal Platform ([GitHub](#))

Tech Stack: dbt, Kafka, Spark, Delta Lake, Airflow, Snowflake, Terraform

- Built two dbt projects with tests, transforming raw claims into governed, BI-ready Snowflake marts
- Built a rolling-baseline fraud detection rule in dbt achieving 100% recall and 67% precision
- Implemented Bronze/Silver/Gold medallion architecture on Delta Lake processing 50,000 synthetic insurance claims

SaaS Subscription Churn & Revenue Analytics Platform ([GitHub](#))

Tech Stack: DuckDB, pandas, PostgreSQL, Grafana

- Derived 15+ metrics tables including NRR, GRR, and ARPA from raw subscription events
- Modeled the MRR waterfall and running balance in SQL using CTEs and window functions
- Built a 21-panel Grafana dashboard provisioned entirely as version-controlled config, not UI clicks

Conditional Generative Data Augmentation for Image Classification ([GitHub](#))

Tech Stack: Python, PyTorch, Conditional GAN, Conditional VAE-GAN, CNN

- Developed CGAN and Conditional VAE-GAN pipelines in PyTorch across MNIST and Fashion-MNIST
- Improved Fashion-MNIST accuracy from 63.6% to 70.8% using only 100 real samples
- Evaluated seven classifier training strategies, finding weighted real-synthetic mixing performed best overall

TECHNICAL SKILLS

- Data Engineering & Orchestration** dbt, Kafka, Spark, Airflow
- Data & SQL** SQL, DuckDB, PostgreSQL, ETL, Data Pipelines, Data Cleaning, Snowflake
- Analytics & BI** Power BI, DAX, Grafana, Data Visualization
- Programming & Tools** Python, pandas, Git, Docker, CI/CD
- Statistics & Methods** Statistical Inference, Experimental Design, Model Evaluation