

SATHWIK H NAIK

| (240) 438-0411 | sathwikhnaik@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#) |

PROFESSIONAL SUMMARY

Data analyst with hands-on SQL, Power BI, and ETL automation experience at a Fortune 50 healthcare organization. Completed M.S. in Data Science at UMD in 2026, backed by a self-directed analytics capstone across DuckDB, pandas, PostgreSQL, and Grafana.

EDUCATION

University of Maryland, College Park M.S. in Data Science Coursework Statistical Inference, Deep Learning, Data Modeling, Data Visualization, Experimental Design, Advanced Time Series	2024 – 2026 GPA 3.81/4.0
Reva University, India B.Tech, Computer Science and Engineering	2019 – 2023 GPA 3.55/4.0

WORK EXPERIENCE

Elevance Health Software Engineer	Aug 2023 – Aug 2024
- Built and maintained 10+ Power BI dashboards with DAX measures for multiple internal teams - Designed reusable KPI structures, translating operational requirements into consistent, self-serve cross-team metrics - Automated recurring data-preparation and reporting workflows, eliminating 20+ hours of manual weekly effort - Developed ETL workflows to source, clean, and validate healthcare data, improving quality for stakeholder reports - Documented report logic, metric definitions, and validation rules enabling governed self-service reporting	

PROJECTS

SaaS Subscription Churn & Revenue Analytics Platform ([GitHub](#))

Tech Stack: DuckDB, pandas, PostgreSQL, Grafana, Plotly, scikit-learn

- Built a 21-panel Grafana dashboard provisioned entirely as version-controlled config, not UI clicks
- Designed a four-tier dashboard hierarchy with live template-variable filters for plan and channel
- Derived 15+ metrics tables including NRR, GRR, and ARPA from raw subscription events

Healthcare Insurance Claims Analytics & Fraud Signal Platform ([GitHub](#))

Tech Stack: Kafka, Spark, Delta Lake, dbt, Airflow, Snowflake, Terraform

- Built two dbt projects with tests, transforming raw claims into governed, BI-ready Snowflake marts
- Built a rolling-baseline fraud detection rule in dbt achieving 100% recall and 67% precision
- Implemented Bronze/Silver/Gold medallion architecture on Delta Lake processing 50,000 synthetic insurance claims

Conditional Generative Data Augmentation for Image Classification ([GitHub](#))

Tech Stack: Python, PyTorch, Conditional GAN, Conditional VAE-GAN, CNN

- Developed CGAN and Conditional VAE-GAN pipelines in PyTorch across MNIST and Fashion-MNIST
- Improved Fashion-MNIST accuracy from 63.6% to 70.8% using only 100 real samples
- Evaluated seven classifier training strategies, finding weighted real-synthetic mixing performed best overall

TECHNICAL SKILLS

- Analytics & BI** Power BI, DAX, Tableau, Grafana, Plotly, Data Visualization
- Data & SQL** SQL, DuckDB, PostgreSQL, ETL, Data Pipelines, Data Cleaning, Snowflake
- Statistics & Methods** Statistical Inference, Experimental Design, A/B Testing, Time-Series Forecasting
- Programming & Tools** Python, pandas, R, Git, Docker, CI/CD
- ML & AI** PyTorch, scikit-learn, Gradient Boosting, Model Evaluation