

SATHWIK H NAIK

| (240) 438-0411 | sathwikhnaik@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#) |

PROFESSIONAL SUMMARY

Data engineer with hands-on SQL, ETL, and pipeline automation experience at a Fortune 50 healthcare organization. Completed M.S. in Data Science at UMD in 2026, backed by a self-directed capstone across Kafka, Spark, dbt, Airflow, Snowflake, and Terraform.

EDUCATION

University of Maryland, College Park M.S. in Data Science Coursework Statistical Inference, Deep Learning, Data Modeling, Data Visualization, Experimental Design, Advanced Time Series	2024 – 2026 GPA 3.81/4.0
Reva University, India B.Tech, Computer Science and Engineering	2019 – 2023 GPA 3.55/4.0

WORK EXPERIENCE

Elevance Health <i>Software Engineer</i> - Developed ETL workflows to source, clean, and validate healthcare data, improving quality for stakeholder reports - Automated recurring data-preparation and reporting workflows, eliminating 20+ hours of manual weekly effort - Designed reusable KPI structures, translating operational requirements into consistent, self-serve cross-team metrics - Built and maintained 10+ Power BI dashboards with DAX measures for multiple internal teams - Documented report logic, metric definitions, and validation rules enabling governed self-service reporting	Aug 2023 – Aug 2024
--	---------------------

PROJECTS

Healthcare Insurance Claims Analytics & Fraud Signal Platform (GitHub) Tech Stack: Kafka (Redpanda), Spark, Delta Lake/Databricks, dbt, Airflow, Snowflake, Terraform - Built a hybrid lakehouse-warehouse pipeline using Kafka, Spark, Delta Lake, dbt, Airflow, and Snowflake - Orchestrated the end-to-end pipeline in Airflow with retries, exponential backoff, and failure alerting - Provisioned Databricks and Snowflake infrastructure declaratively using Terraform with key-pair JWT authentication
SaaS Subscription Churn & Revenue Analytics Platform (GitHub) Tech Stack: DuckDB, pandas, PostgreSQL - Modeled the MRR waterfall and running balance in SQL using CTEs and window functions - Computed cohort retention matrices using CROSS JOIN and FILTER-based conditional aggregation in SQL - Derived 15+ metrics tables including NRR, GRR, and ARPA from raw subscription events
Banking Liquidity Forecast Lab (GitHub) Tech Stack: Python, LightGBM, FastAPI, Docker Compose - Served the forecasting model through FastAPI endpoints packaged in reproducible Docker Compose containers - Engineered leak-safe lag, rolling-window, and Fourier calendar features for multi-segment LightGBM forecasting - Achieved 9.48% WAPE forecasting daily cash outflow with Holt-Winters under six-fold walk-forward validation

TECHNICAL SKILLS

- Data Engineering & Orchestration** dbt, Kafka, Spark, Airflow
- Data & SQL** SQL, ETL, Data Pipelines, Data Cleaning, Snowflake, DuckDB, PostgreSQL
- Programming & Tools** Python, pandas, R, Git, Docker, FastAPI, CI/CD
- Statistics & Methods** Statistical Inference, Time-Series Forecasting, Experimental Design
- ML & AI** PyTorch, scikit-learn, Model Evaluation